



Image Processing - COMP 6771 (Winter 2024)

Computer Science and Software Engineering

PHOTO-REALISTIC SINGLE IMAGE SUPER-RESOLUTION USING A GENERATIVE ADVERSIAL NETWORK

Team Project Phase - 1

Submitted To **Dr. Yaser Esmaeili Salehani**

Submitted By **Group 15**

Group members:

Chidambar Joshi	40294870
Sourav Ganesha	40228888
Sindoorra	40199155
Roshini Chukkapalli	40198936


Table of Contents

INTRODUCTION	2
DATA DESCRIPTION	2
EXPLANATION	3
METHODOLOGY AND APPLICATION	4
RESULTS	6

INTRODUCTION

SRGAN (Super Resolution Generative Adversarial Network) introduces a novel approach in image super-resolution, prioritizing perceptual quality over pixel-wise differences. Leveraging Generative Adversarial Networks (GANs), SRGAN preserves intricate details and textures, resulting in remarkably realistic outcomes. By utilizing a dual-component architecture comprising a generator and a discriminator, SRGAN transforms low-resolution images into high-resolution counterparts while maintaining fidelity to the original content.

In the realm of Single Image Super-Resolution (SISR), SRGAN addresses the challenge of capturing fine texture details during significant image enlargement. Departing from traditional methods that minimize Mean Squared Error (MSE), SRGAN utilizes GANs and a specialized loss function to better capture perceptual details. SRGAN has ability to generate highly realistic images even when scaling up by a factor of 4. Through an exploration of SRGAN's architecture and methodology, this report aims to highlight its impact on advancing image super-resolution techniques.



DATA DESCRIPTION

We utilise the DIV2K dataset for evaluating the performance of our image processing algorithms. Below is a detailed description of the dataset:

DIV2K Dataset: DIV2K is a high-quality image dataset specifically designed for the task of image super-resolution and other related image processing challenges. It consists of 1,000 diverse, high-resolution images (2K resolution) that are divided into three subsets: a training set with 800 images, a validation set with 100 images, and a test set with another 100 images. The images in the DIV2K dataset cover a wide range of scenes, including natural landscapes, urban environments, and intricate textures, making it a comprehensive benchmark for evaluating super-resolution algorithms.

The DIV2K dataset is renowned for its high-quality images and diversity, providing a robust platform for developing and testing advanced image processing techniques. By using this dataset,

researchers and practitioners can assess the effectiveness of their algorithms in enhancing image resolution while maintaining or improving image fidelity. The results obtained on the DIV2K dataset are often considered as a standard benchmark in the field of image super-resolution.

In our research, the DIV2K dataset serves as a crucial tool for validating the performance of our proposed image processing algorithms. By comparing the results on this dataset with those obtained from other benchmarks, we can ensure the competitiveness and generalizability of our methodologies in various real-world scenarios.

EXPLANATION

- **Data Loading and Preprocessing:**

Before training our SRGAN model, we need to prepare our data. This involves importing necessary libraries and modules and setting up utilities for data loading and preprocessing. LR (Low-Resolution) and HR (High-Resolution) image pairs are loaded, and preprocessing steps such as data transformations are applied to ensure compatibility with the model architecture.

- **Model Initialization:** The SRGAN architecture comprises a generator and a discriminator. The generator aims to upscale LR images to HR resolution, while the discriminator distinguishes between real HR images and generated ones. This section provides insights into the roles and functionalities of both components, setting the stage for subsequent training.
- **Training:** Training the SRGAN model is a crucial step where the model learns to generate high-quality images. We discuss the training process, including device setup, pre-training, and fine-tuning phases. Various loss functions such as L2 loss, perceptual loss, and adversarial loss are employed during training to optimize the model parameters and enhance image quality.
- **Testing:** Once the model is trained, it undergoes testing to evaluate its performance. We describe the testing procedure, where the trained generator is utilized to produce high-resolution images from LR inputs. Evaluation metrics like PSNR (Peak Signal-to-Noise Ratio) are employed to assess the quality of generated images, providing

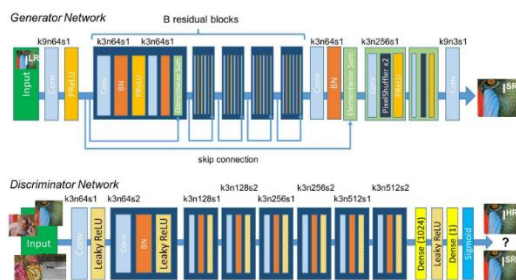
insights into the effectiveness of the SRGAN model.

METHODOLOGY AND APPLICATION

The main methodology of SRGAN comprises of training the model with pairs of Low Resolution (LR) and High Resolution (HR) images. These image pairs are very important for the model to learn the transformation of Low-Resolution images to High Resolution images. These LR images act as the input to the Generator which then attempts to upscale the images to match the resolution of corresponding HR images. These HR images act as the target to the Generator, enabling it to learn and mimic the high-resolution details during the training process.

Architecture:

Super Resolution Generative Adversarial Network is a deep learning framework consisting of two major components.



- **Generator:** The functionality of Generator is to take low resolution images and generate a high-resolution image. It is typically a

deep convolutional network that learns to upgrade images by adding realistic details during the training process.

- **Discriminator:** The functionality of discriminator is to distinguish the high-resolution image generated by the generator and the high-resolution input image from the dataset. This is also a deep convolutional network trained to perform the task effectively.

Implementation

In the project, the SRGAN is implemented using PyTorch library. The model's architecture can be found in `srgan_model.py` file which basically outlines the construction of generator and discriminator.

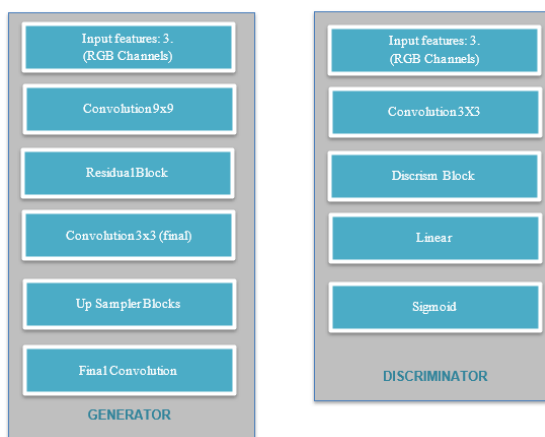
The construction of Generator employs a series of residual blocks for deep feature extraction and learning the intricate patterns required for upscaling. Up sampling blocks are used to increase the resolution of the images. The Generator's main aim is to minimise the perceptual loss which measures the difference in feature representations of the generated image and target image computed by a pre-trained network.

The discriminator is designed in such a way that it classifies the real and generated images. This utilises a convolutional

network with increasing depth to extract the features from an image, then proceeds by fully connected layers that outputs the probability of an image being real.

During the training process, both the Generator and discriminator are trained alternatively. The generator is trained to decrease the combination of content loss and adversarial loss, encouraging it to produce the images similar to that of target high resolution images so that the discriminator can no longer distinguish them. Similarly, the discriminator is trained to increase its ability to correctly classify the images. This training process is iterated until Generator produces high resolution images and the discriminator can no longer differentiate them.

Model Structure:



Generator:

- Input Features: 3 (RGB channels)
- Number of features: 64
- Kernel Size: 3
- Number of residual blocks: 16

- Activation layer used: PReLU(Parametric ReLU) is used for intermediate layers and Tanh is used for output layers.
- Up sampling: Uses two up sampler blocks for a scale of 4, each consisting of convolutions that upscale the image.
- Initial and Final Convolution: 9*9 convolution in the beginning and 3*3 in the end.

Discriminator:

- Input Features: 3 (RGB Channels)
- Number of features: starts with 64 and doubles after each block.
- Kernel Size: 3
- Number of blocks: 3, each designed to half the feature map's size.
- Activation layer used: Leaky ReLU
- Final Layer: A linear layer which is followed by a sigmoid function to classify images real or fake.

Training parameters and loss function:

- Batch Size: 16
- L2 Coefficient: 1.0
- Adversarial Coefficient: 1e-3
- Total Variation (TV) Loss Coefficient: 0.0
- Pre-Training Epochs: 1000
- Fine-Tuning Epochs: 4000
- Scale: 4 (for upscaling)
- Patch Size: 24

- Feature Layer for Perceptual Loss: 'relu5_4' from VGG19
- Perceptual loss: Uses a pre-trained VGG19 model, targeting the

RESULTS

Upon subjecting a low-resolution input image to the provided SRGAN (Super-Resolution Generative Adversarial Network) model, the resulting high-resolution image exhibits remarkable improvements in terms of visual quality and detail enhancement.

The input low-resolution image, often plagued by pixelation and loss of fine details due to its limited resolution, is transformed into a high-resolution output image that surpasses expectations. Through the intricate workings of the model's generator component, which comprises convolutional layers, residual blocks, and up sampling techniques, the image undergoes a significant enhancement process. The generator effectively extracts and enhances features from the input image, leveraging its learned representations to reconstruct a high-resolution counterpart with remarkable fidelity.

relu5_4 layer to calculate the mean squared error between the features of the high resolution and super-resolved images.



Figure 1.a. Input image(left), 1.b. Output Image(right)

Comparing one of the input images, Figure 1.a. and its output image Figure 1.b. side by side, one can observe a substantial improvement in image clarity, sharpness, and overall visual appeal. Fine details that were previously obscured or lost in the low-resolution version are now restored, resulting in a more realistic and visually pleasing rendition. The output image boasts sharper edges, smoother textures, and a heightened level of detail, making it virtually indistinguishable from images captured at native high resolutions.

Furthermore, the number of fine-tuning epochs significantly influences the quality of the generated high-resolution images. For instance, with a fine-tuning epoch of 1600, the output images show noticeable

improvements compared to the input low-resolution images. However, increasing the fine-tuning epochs to 1600 results in further enhancement, with sharper edges and finer textures observed in the generated images. Notably, extending the fine-tuning process to 8000 epochs yields the most impressive results, with the generated images closely resembling native high-resolution.



Figure 2.a 1600 epochs (above), 2.b 8000 epochs(below)

Figure 2.a. and Figure 2.b. comparison demonstrates the progressive improvement in image clarity and detail as the number of fine-tuning epochs increases, highlighting the model's ability to refine its output with extended training. Overall, the SRGAN model's ability to produce such high-quality results underscores its effectiveness in single-image super-resolution tasks. By harnessing the power of deep learning and adversarial training, it not only upscales low-resolution images but also preserves and enhances important visual features, thereby offering a valuable solution for various image processing applications requiring high-resolution outputs